

# Support for Cloud Native Storage (CNS) on TKGI Clusters with Stateful Application Example

Support for Cloud Native Storage (CNS) on TKGI Clusters with Stateful Application Example.....	1
Introduction.....	1
Overview.....	1
Prerequisites for CNS with TGKI.....	3
Install CSI Driver on a TGKI Cluster .....	3
Deploy Stateful Containerized Application .....	8
Install Velero Backup/Restore product w/support for CSI volume snapshots .....	12
Backup Stateful App with vSphere Volume Snapshot .....	14
Delete & Restore Stateful App from Backup with vSphere Volume Snapshot.....	15
Troubleshooting Notes.....	17
Conclusion .....	19

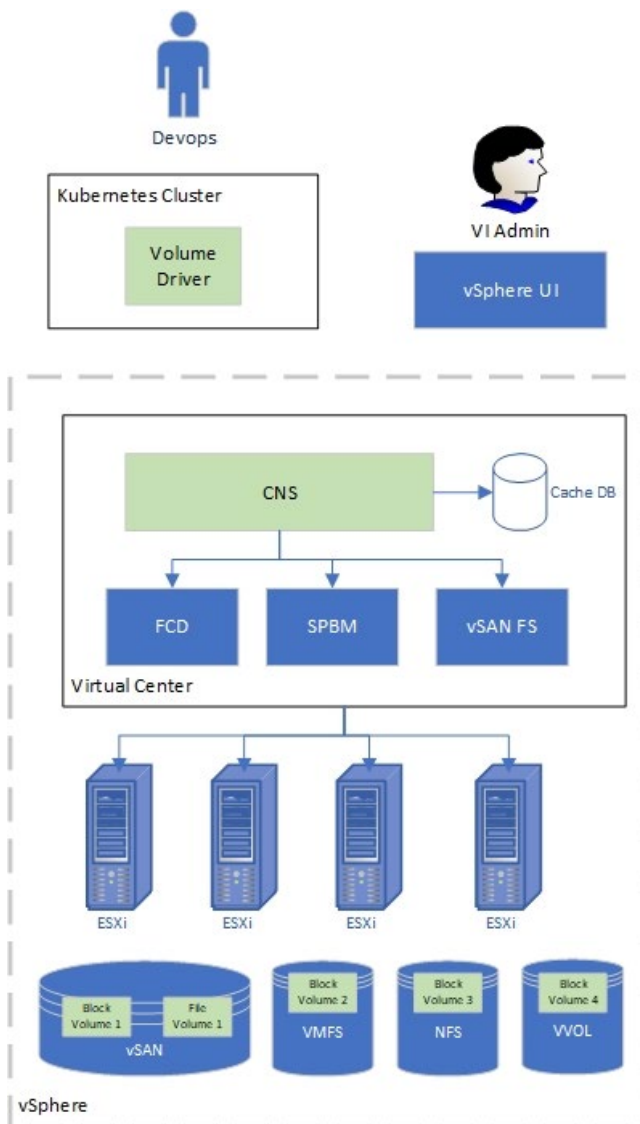
## Introduction

This document is a quick start guide for enabling CNS volume support on a Tanzu Kubernetes Grid Integrated (TKGI, formerly known as Enterprise PKS) Kubernetes clusters and deploying applications using persistent volumes. This document will provide details on configuration of CSI drivers in designated K8s cluster, deployment a stateful application using K8s storage class with CSI driver (OPTIONAL: and backing it up/restoring using Velero open source solution.

Note: thanks to [Alexander Ullah](#) for creating a very helpful "how to" [blog](#) for workload migration between K8s clusters with CSI which I took definition of the stateful [Ghost application](#) from.

## Overview

Cloud Native Storage (CNS) provides comprehensive data management for stateful, containerized apps, enabling apps to survive restarts and outages.



CNS vSphere offers the following two components:

- CNS in vCenter Server
- vSphere volume driver in Kubernetes

CNS control plane introduces a concept of volumes: **container volumes and persistent volumes** in vSphere - it is the Storage control plane for container volumes. CNS is responsible for managing the lifecycle of volumes, including such operations as create, read, update, and delete.

It is also responsible for managing volume metadata, snapshot and restore, volume copy and clone, as well as monitoring the health and compliance of volumes.

These volumes are independent of the virtual machine lifecycle and have their own identity in vSphere.

CNS supports block volumes backed by First Class Disk (FCD) and file volumes backed by vSAN file shares.

A **block volume** can only be attached to one Kubernetes pod with **ReadWriteOnce** access mode at any point in time. A **file volume** can be attached to one or more pods with **ReadWriteMany/ReadOnlyMany** access modes.

In Kubernetes, CNS provides a volume driver that has two sub-components: the **CSI driver** and the **syncer**.

- The **CSI driver** is responsible for volume provisioning, attaching and detaching the volume to VMs, mounting, formatting and unmounting volumes from the pod within the node VM, etc. The CSI driver is built as an out-of-tree CSI plugin for Kubernetes.
- The **syncer** is responsible for pushing PV, PVC, and pod metadata to CNS. It also has a CNS operator that is used in the context of vSphere with Kubernetes (formerly “Project Pacific”).

Stateful containers can use vSphere storage primitives - standard volume, persistent volume, and dynamic provisioning - independent of VM and container lifecycle. vSphere storage backs the volumes, and we can set a storage policy directly on the volumes.

After creation of volumes, we can review them and their backing virtual disks, and monitor their storage policy compliance via use vSphere client.

## Prerequisites for CNS with TGKI

(See <https://docs.pivotal.io/pks/1-7/vsphere-cns.html> for complete list of pre-requisites)

## Install CSI Driver on a TGKI Cluster

Follow documentation: <https://docs.pivotal.io/pks/1-7/vsphere-cns.html> and use sample files provided

(provided input to Docs team for some inaccuracies which were corrected)

1. Create CSI Secret based on provided sample deployment file (*csi-vsphere.conf*)

```
[Global]

cluster-id = PKS-COMPUTE-EDGES

[VirtualCenter "192.168.2.20"]

insecure-flag = "true"

user = "administrator@vsphere.local"

password = "XXXXXX"

port = "443"

datacenters = "PKS-NESTED-DC2"
```

### NOTES:

*cluster-id* is a unique identifier, can be a name of vSphere cluster which data store will be used  
*VirtualCenter* contains IP address (FQDN) of vSphere vCenter  
*user* and *password* are credentials of vSphere admin user (or user with sufficient access level)  
*datacenters* are vSphere Data center(s) where clusters are contained

Create a secret for accessing vSphere Data Center/Cluster

```
kubectl create secret generic vsphere-config-secret --from-file=csi-vsphere.conf --  
namespace=kube-system  
secret/vsphere-config-secret created
```

2. Create RBAC objects for CSI access based on provided sample deployment file (*vsphere-csi-controller-rbac.yaml* included into 'samples' repository folder, partially shown below)

```
kind: ServiceAccount  
  
apiVersion: v1  
  
metadata:  
  
  name: vsphere-csi-controller  
  
  namespace: kube-system  
  
---  
  
kind: ClusterRole  
  
apiVersion: rbac.authorization.k8s.io/v1  
  
metadata:  
  
  name: vsphere-csi-controller-role  
  
rules:  
  
  - apiGroups: [""]  
  
    resources: ["nodes", "persistentvolumeclaims", "pods"]  
  
    verbs: ["get", "list", "watch"]  
  
...
```

```
kubectl apply -f vsphere-csi-controller-rbac.yaml
```

```
serviceaccount/vsphere-csi-controller created  
clusterrole.rbac.authorization.k8s.io/vsphere-csi-controller-role created  
clusterrolebinding.rbac.authorization.k8s.io/vsphere-csi-controller-binding created
```

Verify that service account, cluster roles/cluster role binding exist in 'kube-system' namespace:

```
kubectl get serviceaccounts -n kube-system
```

```
.....  
  
vrealize-med01          1      25h  
  
vsphere-csi-controller    1      3m4s
```

## kubectl get clusterroles -n kube-system

vrops-cadvisor	47h
vsphere-csi-controller-role	<invalid>

## kubectl get clusterrolebindings -n kube-system

vrops-cadvisor	47h
vsphere-csi-controller-binding	<invalid>

3. Install the vSphere CSI Driver using sample manifest file (*vsphere-csi-controller-ss.yaml* included into 'samples' repository folder)

## kubectl apply -f vsphere-csi-controller-ss.yaml

statefulset.apps/vsphere-csi-controller created

csidriver.storage.k8s.io/csi.vsphere.vmware.com created

Optionally, monitor events in kube-system namespace to check initialization of vsphere-csi-controller:

## kubectl get events -n kube-system

LAST SEEN	TYPE	REASON	OBJECT	MESSAGE
<unknown>	Normal	Scheduled	pod/vsphere-csi-controller-0	Successfully assigned kube-system/vsphere-csi-controller-0 to a208e18a-fbe4-4e36-a17f-6efe7c204d78
27s	Normal	Pulling	pod/vsphere-csi-controller-0	Pulling image " <a href="https://quay.io/k8scsi/csi-attacher:v1.1.1">quay.io/k8scsi/csi-attacher:v1.1.1</a> "
22s	Normal	Pulled	pod/vsphere-csi-controller-0	Successfully pulled image " <a href="https://quay.io/k8scsi/csi-attacher:v1.1.1">quay.io/k8scsi/csi-attacher:v1.1.1</a> "
20s	Normal	Created	pod/vsphere-csi-controller-0	Created container csi-attacher
20s	Normal	Started	pod/vsphere-csi-controller-0	Started container csi-attacher
20s	Normal	Pulling	pod/vsphere-csi-controller-0	Pulling image " <a href="https://gcr.io/cloud-provider-vsphere/csi/release/driver:v1.0.2">gcr.io/cloud-provider-vsphere/csi/release/driver:v1.0.2</a> "
7s	Normal	Pulled	pod/vsphere-csi-controller-0	Successfully pulled image " <a href="https://gcr.io/cloud-provider-vsphere/csi/release/driver:v1.0.2">gcr.io/cloud-provider-vsphere/csi/release/driver:v1.0.2</a> "
3s	Normal	Created	pod/vsphere-csi-controller-0	Created container vsphere-csi-controller
2s	Normal	Started	pod/vsphere-csi-controller-0	Started container vsphere-csi-controller
2s	Normal	Pulling	pod/vsphere-csi-controller-0	Pulling image " <a href="https://quay.io/k8scsi/livenessprobe:v1.1.0">quay.io/k8scsi/livenessprobe:v1.1.0</a> "
30s	Normal	SuccessfulCreate	statefulset/vsphere-csi-controller	create Pod vsphere-csi-controller-0 in StatefulSet vsphere-csi-controller successful

4. Install CSI Driver DaemonSet using sample manifest file (*vsphere-csi-node-ds.yaml* included into 'samples' repository folder)

**kubectl apply -f vsphere-csi-node-ds.yaml**

daemonset.apps/vsphere-csi-node created

Verify that DaemonSet (that is guaranteed Pod/Node) is running on all cluster nodes:

**kubectl get ds -n kube-system**

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
vsphere-csi-node	4	4	4	4	4	<none>	2m8s

## 5. Verify that CSI Driver Deployed Successfully

Check that all pods are running in the kube-system namespace (look for **vsphere-csi-...** pod names)

**kubectl get po --namespace=kube-system**

NAME	READY	STATUS	RESTARTS	AGE
------	-------	--------	----------	-----

NAME	READY	STATUS	RESTARTS	AGE
------	-------	--------	----------	-----

coredns-5b6649768f-7zzfr	1/1	Running	0	97m
--------------------------	-----	---------	---	-----

coredns-5b6649768f-8rg9m	1/1	Running	0	97m
--------------------------	-----	---------	---	-----

coredns-5b6649768f-c4ktq	1/1	Running	0	97m
--------------------------	-----	---------	---	-----

metrics-server-7f9887fbb5-64rtk	1/1	Running	0	97m
---------------------------------	-----	---------	---	-----

<b>vsphere-csi-controller-0</b>	<b>5/5</b>	<b>Running</b>	<b>0</b>	<b>7m</b>
---------------------------------	------------	----------------	----------	-----------

<b>vsphere-csi-node-b47k4</b>	<b>3/3</b>	<b>Running</b>	<b>0</b>	<b>5m45s</b>
-------------------------------	------------	----------------	----------	--------------

<b>vsphere-csi-node-cvcm2</b>	<b>3/3</b>	<b>Running</b>	<b>0</b>	<b>5m45s</b>
-------------------------------	------------	----------------	----------	--------------

<b>vsphere-csi-node-fxfsj</b>	<b>3/3</b>	<b>Running</b>	<b>0</b>	<b>5m45s</b>
-------------------------------	------------	----------------	----------	--------------

<b>vsphere-csi-node-p4wgc</b>	<b>3/3</b>	<b>Running</b>	<b>0</b>	<b>5m45s</b>
-------------------------------	------------	----------------	----------	--------------

## 6. Verify that CRDs for CSI are deployed and working:

**kubectl get CSINode**

NAME	CREATED AT
23f8a803-50de-4d2a-be42-5f4482f35fb3	2020-06-25T19:31:44Z

3f178b02-514c-415d-9716-470b30175b44 2020-06-30T05:11:41Z  
576e0ba9-5621-4ecd-a1b1-eb39af781390 2020-06-25T19:31:42Z  
9dcd874c-c86e-4b59-af77-c914619fd38c 2020-06-10T19:23:29Z

### **kubectl describe CSINode**

Name: 23f8a803-50de-4d2a-be42-5f4482f35fb3  
Namespace:  
Labels: <none>  
Annotations: <none>  
API Version: [storage.k8s.io/v1beta1](https://storage.k8s.io/v1beta1)  
Kind: CSINode  
Metadata:  
Creation Timestamp: 2020-06-25T19:31:44Z  
Owner References:  
API Version: v1  
Kind: Node  
Name: 23f8a803-50de-4d2a-be42-5f4482f35fb3  
UID: 2e495151-1a4a-4aac-baab-f5182b504e21  
Resource Version: 7801510  
Self Link: /apis/[storage.k8s.io/v1beta1](https://storage.k8s.io/v1beta1)/csinodes/[23f8a803-50de-4d2a-be42-5f4482f35fb3](https://storage.k8s.io/v1beta1)  
UID: efl1d6c57-3965-499b-8493-5952e454dadc  
Spec:  
Drivers:  
Name: [csi.vsphere.vmware.com](https://csi.vsphere.vmware.com)  
Node ID: 23f8a803-50de-4d2a-be42-5f4482f35fb3  
Topology Keys: <nil>  
Events: <none>

....

### 7. Verification that CSI Driver is installed

#### **kubectl get csidrivers**

NAME	CREATED AT
csi.vsphere.vmware.com	2020-07-03T02:42:08Z

### 8. Verify that CSI ProviderID was added to all Cluster Nodes

#### **kubectl describe nodes | grep "ProviderID" \*\***

ProviderID:	vsphere://421c57b9-99bb-ef5a-c2d4-38b857f8d9d5
ProviderID:	vsphere://421cef46-3327-d66c-4d52-5e6f2bb1bfaa
ProviderID:	vsphere://421c04c4-13fd-069e-c2a9-d13166d81a4c
ProviderID:	vsphere://421ce760-a829-deb2-82d9-45ff4fc5ec95

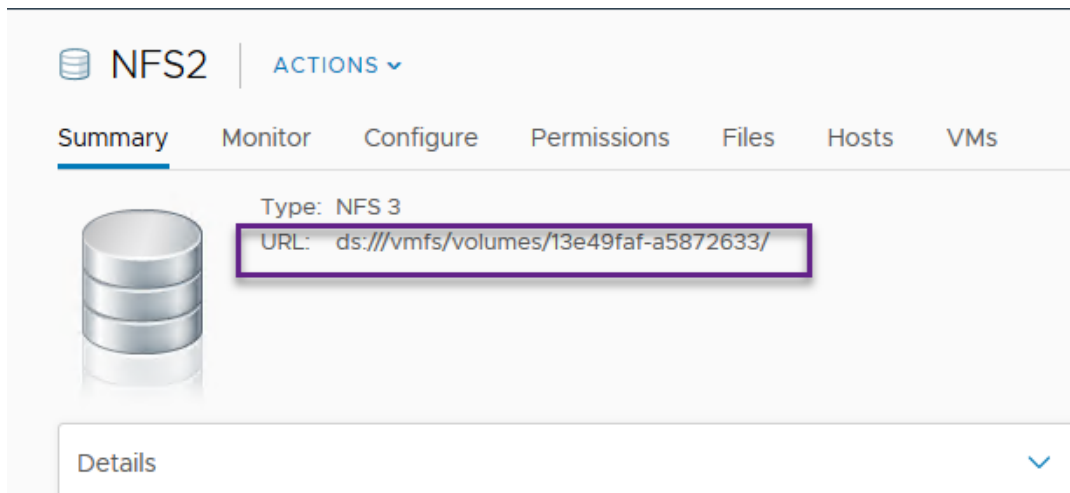
# Deploy Stateful Containerized Application

We will simple simple Ghost blogging application that preserves state using PVC (source code available from [https://github.com/beyondelastic/velero\\_vSphere](https://github.com/beyondelastic/velero_vSphere))

1. Create storage class that is using CSI driver (using csi-sc.yaml deployment descriptor file available in 'samples' folder)

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: demo-sts-sc
  annotations:
    storageclass.kubernetes.io/is-default-class: "true"
provisioner: csi.vsphere.vmware.com
parameters:
  datastoreurl: "ds:///vmfs/volumes/13e49faf-a5872633/"
```

NOTE: datastoreurl should point to a folder in the associated vSphere data store:



**kubectl apply -f csi-sc-dan.yaml -n ghost**

storageclass.storage.k8s.io/demo-sts-sc created

Verify that storage class primitive has been created at K8s cluster level:

**kubectl get sc**

NAME	PROVISIONER	RECLAIMPOLICY	VOLUMEBINDINGMODE	ALLOWVOLUMEEXPANSION	AGE
demo-sts-sc (default)	csi.vsphere.vmware.com	Delete	Immediate	false	4m40s



2. Create PVS using SC above and deploy Ghost app based on that PVC (ghost-claim.yaml sample deployment descriptor is included into the ‘samples’ repository folder)

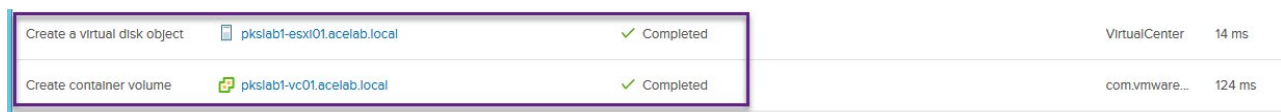
```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: blog-content-new
  namespace: ghost
  annotations:
    volume.beta.kubernetes.io/storage-class: demo-sts-sc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 2Gi
```

Run the following command:

**kubectl apply -f ghost-claim.yaml**

persistentvolumeclaim/blog-content-new created

NOTE: on vSphere UI we can monitor tasks of creation of persistent volume:



Create a virtual disk object	pkslab1-esxi01.acefab.local	✓ Completed	VirtualCenter	14 ms
Create container volume	pkslab1-vc01.acefab.local	✓ Completed	com.vmware...	124 ms

Verify that PVC that volume are in “bound” state:

**kubectl get pvc -n ghost**

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
<b>blog-content-new</b>	<b>Bound</b>	<b>pvc-f9ec1bce-c2c1-4477-b118-e3b333a57151</b>	<b>2Gi</b>	<b>RWO</b>	<b>demo-sts-sc</b>	<b>119s</b>

3. Deploy Ghost stateful application (using ghost-new.yaml sample deployment descriptor file found in the ‘samples’ folder, shown below):

```
apiVersion: v1

kind: Service

metadata:

  labels:

    name: blog

  name: blog

  namespace: ghost
```

```
spec:

  ports:

    - port: 80

      targetPort: 2368

  selector:

    app: blog

  type: LoadBalancer
---

apiVersion: apps/v1

kind: Deployment

metadata:

  name: blog

  namespace: ghost

  labels:

    app: blog

spec:

  replicas: 1

  selector:

    matchLabels:

      app: blog

  template:

    metadata:

      labels:

        app: blog

    spec:

      containers:

        - name: blog

          image: ghost:latest
```

```

    imagePullPolicy: Always

    ports:

    - containerPort: 2368

    env:

    - name: url

      value: http://my-blog.acelab.local ← needs to be registered in DNS

    volumeMounts:

    - mountPath: /var/lib/ghost/content

      name: content

    volumes:

    - name: content

      persistentVolumeClaim:

        claimName: blog-content-new ← needs to match PVC name created in the
previous step

```

-----

**kubectl apply -f ghost-new.yaml**

service/blog created

deployment.apps/blog created

4. Verify that all application objects have been created in the designated K8s namespace:

**kubectl get all -n ghost**

NAME	READY	STATUS	RESTARTS	AGE
pod/blog-8646cd4d56-fpqz9	1/1	Running	0	10m


NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/blog	LoadBalancer	10.100.200.102	192.168.74.81	80:35713/TCP	10m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/blog	1/1	1	1	10m

.....

Navigate to the URL defined by EXTERNAL-IP address of LoadBalancer service:



GETTING STARTED

## Welcome to Ghost

Welcome, it's great to have you here. We know that first impressions are important, so we've populated your new site with some initial getting started posts that will help you get familiar with everything in no time.

 GHOST  
1.0.0-rc.2020-08-14 14:11:00

Now we can start using the Blog application as it is intended – posting blogs etc.

## Install Velero Backup/Restore product w/support for CSI volume snapshots

1. Install minio for local object store using [Helm3 chart](#) from VMware Bitnami:

```
helm install minio-release -n minio \
--set access.Key.password=minio \
--set secretKey.password=minio123 \
--set persistence.storageClass=thin-disk \
bitnami/minio
```

NOTE: If access.Key.password and secretKey.password values are different from those passed into Helm chart, you can retrieve their values by:

```
export MINIO_ACCESS_KEY=$(kubectl get secret --namespace minio minio-release -o
jsonpath="{.data.access-key}" | base64 --decode)
```

```
export MINIO_SECRET_KEY=$(kubectl get secret --namespace minio minio-release -o
jsonpath="{.data.secret-key}" | base64 --decode)
```

use outputs of those commands in **credentials** file to provide access to S3 object store (minIO) used in the script below:

[default]

aws\_access\_key\_id=mtTK1JFdwQ

aws\_secret\_access\_key=minio123

2. Use the following command script to install Velero w/support for CSI:

```
export BUCKET=velero
export REGION=minio
./velero install \
  --provider aws \
  --bucket $BUCKET \
  --secret-file ./credentials \
  --backup-location-config
region=$REGION,s3ForcePathStyle="true",s3Url=http://192.168.74.104 \
  --snapshot-location-config region=$REGION \
  --plugins velero/velero-plugin-for-aws:v1.1.0
```

Follow <https://github.com/vmware-tanzu/velero-plugin-for-vsphere> to install vSphere snapshot plugin (v 1.0.1!)

3. Add vSphere volume snapshot plugin

**velero plugin add vsphereveleroplugin/velero-plugin-for-vsphere:1.0.1**

and create vSphere volume snapshot location:

**velero snapshot-location create vsl-vsphere-dz --provider [velero.io/vsphere](https://velero.io/vsphere)**

Snapshot volume location "vsl-vsphere-dz" configured successfully.

(Steps below are optional)

verify that snapshot location has been created:

```
velero snapshot-location get
NAME          PROVIDER
default       aws
vsl-vsphere-dz velero.io/vsphere
```

Optionally, verify that all pods, including "datamgr." to handle volume snapshots, are running in the velero namespace:

**kubectrl get po -n velero**

NAME	READY	STATUS	RESTARTS	AGE
pod/datamgr-for-vsphere-plugin-d5hbl	0/1	ContainerCreating	0	4m22s
pod/datamgr-for-vsphere-plugin-lmrpw	1/1	Running	0	4m22s
pod/datamgr-for-vsphere-plugin-mwgcw	1/1	Running	0	4m22s

## Backup Stateful App with vSphere Volume Snapshot

1. Using volume snapshot above, initiate backup of stateful application Ghost namespace configured earlier:

```
velero backup create ghost-backup3 --include-namespaces=ghost --snapshot-volumes --volume-snapshot-locations vsl-vsphere-dz
```

Backup request "ghost-backup1" submitted successfully.

Run `velero backup describe ghost-backup1` or `velero backup logs ghost-backup1` for more details.

2. Verify that backup completed successfully and volume snapshots were included:

```
velero describe backup ghost-backup3 --details
```

```
-----
Name:      ghost-backup3
Namespace: velero
Labels:    velero.io/storage-location=default
Annotations: velero.io/source-cluster-k8s-gitversion=v1.16.7+vmware.1
             velero.io/source-cluster-k8s-major-version=1
             velero.io/source-cluster-k8s-minor-version=16
```

Phase: Completed

Namespaces:

**Included: ghost**

Excluded: <none>

Resources:

Included: \*

Excluded: <none>

Cluster-scoped: auto

Label selector: <none>

Storage Location: default

**Velero-Native Snapshot PVs: true**

TTL: 720h0m0s

Hooks: <none>

Backup Format Version: 1

Started: 2020-07-07 19:19:55 +0000 UTC

Completed: 2020-07-07 19:20:35 +0000 UTC

Expiration: 2020-08-06 19:19:55 +0000 UTC

**Total items to be backed up: 10**

**Items backed up: 10**

Resource List:

apps/v1/Deployment:

- ghost/blog

apps/v1/ReplicaSet:  
- ghost/blog-8646cd4d56  
v1/Endpoints:  
- ghost/blog  
v1/Namespace:  
- ghost  
**v1/PersistentVolume:**  
- **pvc-f9ec1bce-c2c1-4477-b118-e3b333a57151**  
v1/PersistentVolumeClaim:  
- ghost/blog-content-new  
v1/Pod:  
- ghost/blog-8646cd4d56-fpqz9  
v1/Secret:  
- ghost/default-token-c6zpx  
v1/Service:  
- ghost/blog  
v1/ServiceAccount:  
- ghost/default

Velero-Native Snapshots:

**pvc-f9ec1bce-c2c1-4477-b118-e3b333a57151:**

Snapshot ID:      ivd:be575613-2aae-4823-9e6d-16d4ce3a6fca:651985e4-642a-4a78-93e3-976f566d652d

Type:              ivd

Availability Zone:

IOPS:              100

-----  
NOTES:

- notice the **PersistentVolume** object being included into the backup and corresponding Velero Native Snapshots, it should match the PVC ID of app created by deployment (see above)
- uploading volume snapshots may take a long time due to many processes happening on the background, make sure to check backup upload status:

**kubectl -n velero get uploads**

## Delete & Restore Stateful App from Backup with vSphere Volume Snapshot

1. Delete namespace where app is deployed:

**kubectl delete ns ghost**

make sure the application namespace is fully deleted

2. Initiate a restore process using previously created backup (with volume snapshots) :

**velero restore create restore-ghost-backup3 --from-backup ghost-backup3**

Restore request “restore-ghost-backup3” submitted successfully.

Run 'velero restore describe restore-ghost-backup3' or 'velero restore logs restore-ghost-backup3' for more details

3. Monitor restore process status/check logs:

**velero restore describe restore-ghost-backup3**

Name: restore-ghost-backup3

Namespace: velero

Labels: <none>

Annotations: <none>

Phase: Completed

Backup: ghost-backup3

Namespaces:

**Included: all namespaces found in the backup**

Excluded: <none>

Resources:

Included: \*

Excluded: nodes, events, [events.k8s.io](https://events.k8s.io), [backups.velero.io](https://backups.velero.io), [restores.velero.io](https://restores.velero.io), [resticrepositories.velero.io](https://resticrepositories.velero.io)

Cluster-scoped: auto

Namespace mappings: <none>

Label selector: <none>

Restore PVs: auto

-----

4. Check status of restored namespace, PVC, deployments, pods and services:

**kubectrl get ns ghost**

NAME	STATUS	AGE
------	--------	-----

ghost	Active	3m37s
-------	--------	-------

**kubectrl get pvc -n ghost**

NAME	STATUS	VOLUME	CAPACITY	ACCES
S MODES	STORAGECLASS	AGE		

blog-content-new	Bound	pvc-f9ec1bce-c2c1-4477-b118-e3b333a57151	2Gi	RWO	demo-sts-sc	5m41s	<== important to verify PVC is bound
------------------	-------	--	-----	-----	-------------	-------	--------------------------------------

**kubectrl get po,svc -n ghost**

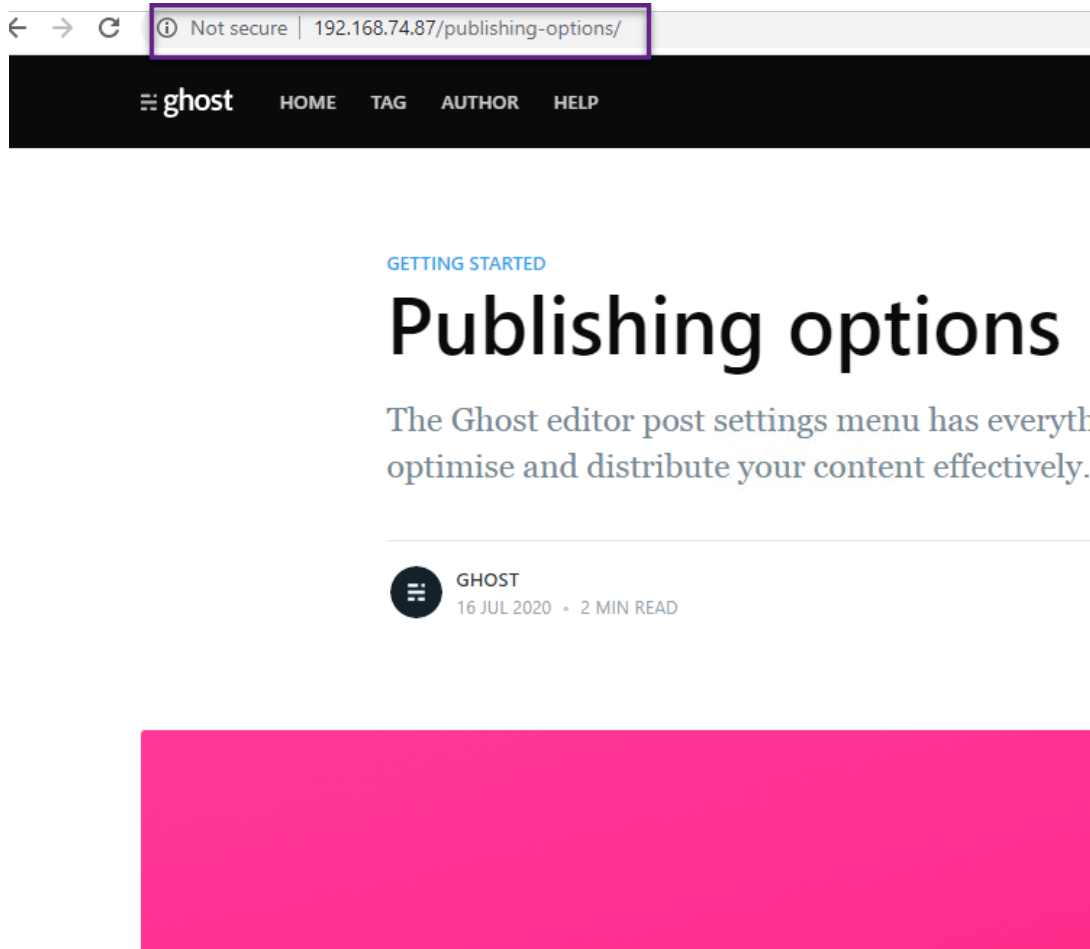
NAME	READY	STATUS	RESTARTS	AGE
------	-------	--------	----------	-----

pod/blog-8646cd4d56-fpqz9	1/1	Running	0	17m
---------------------------	-----	---------	---	-----



NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/blog	LoadBalancer	10.100.200.246	<b>192.168.74.81</b>	80:34241/TCP	17m

Finally, navigate to the LoadBalancer URL of restored version of the Ghost app and check its availability (may need to update the DNS record matching URL specified in deployment descriptor):



## Troubleshooting Notes

1. If validation of all Pods implementing CSI driver is failing (e.g. **vsphere-csi-controller** pod) is failing, examine pod logs for RCA:

**kubectl logs vsphere-csi-controller-0 vsphere-csi-controller -n kube-system**

```
I0702 20:16:00.623537    1 config.go:261] GetCnsconfig called with cfgPath: /etc/cloud/csi-vsphere.conf
I0702 20:16:00.624455    1 config.go:206] Initializing vc server your_vc_ip
I0702 20:16:00.624493    1 controller.go:67] Initializing CNS controller
I0702 20:16:00.624663    1 virtualcentermanager.go:63] Initializing defaultVirtualCenterManager...
I0702 20:16:00.624736    1 virtualcentermanager.go:65] Successfully initialized
defaultVirtualCenterManager
I0702 20:16:00.624750    1 virtualcentermanager.go:107] Successfully registered VC "your_vc_ip"
```

```

I0702 20:16:00.624878    1 manager.go:62] Initializing volume.volumeManager...
I0702 20:16:00.624885    1 manager.go:66] volume.volumeManager initialized
E0702 20:16:00.643785    1 virtualcenter.go:127] Failed to create new client with err: Post
https://your\_vc\_ip:443/sdk: dial tcp: lookup your_vc_ip on 192.168.2.10:53: server misbehaving
E0702 20:16:00.644010    1 virtualcenter.go:238] Failed to create govnomi client with err: Post
https://your\_vc\_ip:443/sdk: dial tcp: lookup your_vc_ip on 192.168.2.10:53: server misbehaving
E0702 20:16:00.644251    1 virtualcenter.go:201] Cannot connect to vCenter with err: Post
https://your\_vc\_ip:443/sdk: dial tcp: lookup your_vc_ip on 192.168.2.10:53: server misbehaving
E0702 20:16:00.644408    1 util.go:41] Failed to connect to VirtualCenter host: "your_vc_ip". err=Post
https://your\_vc\_ip:443/sdk: dial tcp: lookup your_vc_ip on 192.168.2.10:53: server misbehaving
E0702 20:16:00.644575    1 controller.go:92] Failed to get vcenter. err=Post https://your\_vc\_ip:443/sdk:
dial tcp: lookup your_vc_ip on 192.168.2.10:53: server misbehaving
E0702 20:16:00.644617    1 service.go:107] Failed to init controller. Error: Post
https://your\_vc\_ip:443/sdk: dial tcp: lookup your_vc_ip on 192.168.2.10:53: server misbehaving
I0702 20:16:00.644726    1 service.go:88] configured: csi.vsphere.vmware.com with
map[mode:controller]
time="2020-07-02T20:16:00Z" level=info msg="removed sock file"
path=/var/vcap/data/csi/sockets/pluginproxy/csi.sock
time="2020-07-02T20:16:00Z" level=fatal msg="grpc failed" error="Post https://your\_vc\_ip:443/sdk: dial
tcp: lookup your_vc_ip on 192.168.2.10:53: server misbehaving"
-----

```

**Resolution:** make sure that DNS server is running and listening on port 53, restart if needed

## 2. Running backup using vSphere volume snapshot:

### velero snapshot-location get

```

NAME      PROVIDER
default    aws
*vsl-vsphere-dz  velero.io/vsphere

```

initiate velero backup using snapshot location:

```

velero backup create ghost-backup2 --include-namespaces=ghost --snapshot-volumes --volume-
snapshot-locations vsl-vsphere-dz

```

Backup request "ghost-backup2" submitted successfully.

Run `velero backup describe ghost-backup2` or `velero backup logs ghost-backup2` for more details.

---

However, when running backup of stateful up using CSI volumes, I am getting the following exception:

```

time="2020-07-07T16:26:21Z" level=error msg="Error getting volume snapshotter for volume snapshot
location" backup=velero/ghost-backup2 error="unable to locate VolumeSnapshotter plugin
named velero.io/vsphere" logSource="pkg/backup/item_backupper.go:437" name=pvc-f9ec1bce-c2c1-
4477-b118-e3b333a57151 namespace=persistentVolume=pvc-f9ec1bce-c2c1-4477-b118-e3b333a57151
resource=persistentvolumes volumeSnapshotLocation=vsl-vsphere-dz

```

**Resolution:** (worked after a few retries)

## 3. Examine events when restoring application from backup:

### kubectrl get events -n ghost

LAST SEEN	TYPE	REASON	OBJECT	MESSAGE
-----------	------	--------	--------	---------

```
<unknown> Normal Scheduled pod/blog-8646cd4d56-fpqz9 Successfully assigned
ghost/blog-8646cd4d56-fpqz9 to 576e0ba9-5621-4ecd-a1b1-eb39af781390
14m Normal SuccessfulAttachVolume pod/blog-8646cd4d56-fpqz9 AttachVolume.Attach
succeeded for volume "pvc-f9ec1bce-c2c1-4477-b118-e3b333a57151"
79s Warning FailedMount pod/blog-8646cd4d56-fpqz9 MountVolume.MountDevice failed
for volume "pvc-f9ec1bce-c2c1-4477-b118-e3b333a57151" : rpc error: code = Unavailable desc = all
SubConns are in TransientFailure, latest connection error: connection error: desc = "transport: Error while
dialing dial unix /var/vcap/data/kubelet/plugins_registry/csi.vsphere.vmware.com/csi.sock: connect: no
such file or directory"
78s Warning FailedMount pod/blog-8646cd4d56-fpqz9 Unable to attach or mount
volumes: unmounted volumes=[content], unattached volumes=[default-token-cmqfq content]: timed out
waiting for the condition
3m32s Warning FailedMount pod/blog-8646cd4d56-fpqz9 Unable to attach or mount
volumes: unmounted volumes=[content], unattached volumes=[content default-token-cmqfq]: timed out
waiting for the condition
```

**Resolution:** make sure all nodes in K8s cluster are running and pods can be scheduled on them, wait for scheduler to retry (??)

## Conclusion

We hope this document was useful. As you try these configuration steps, please provide any feedback or questions in the comments section for this document on [code.vmware.com](https://code.vmware.com). Also, please let us know if you have any suggestions or if you would like to see guidance on other topics